



Apache Flink™ Training

Advanced Stream Processing

Tzu-Li (Gordon) Tai

tzulitai@apache.org

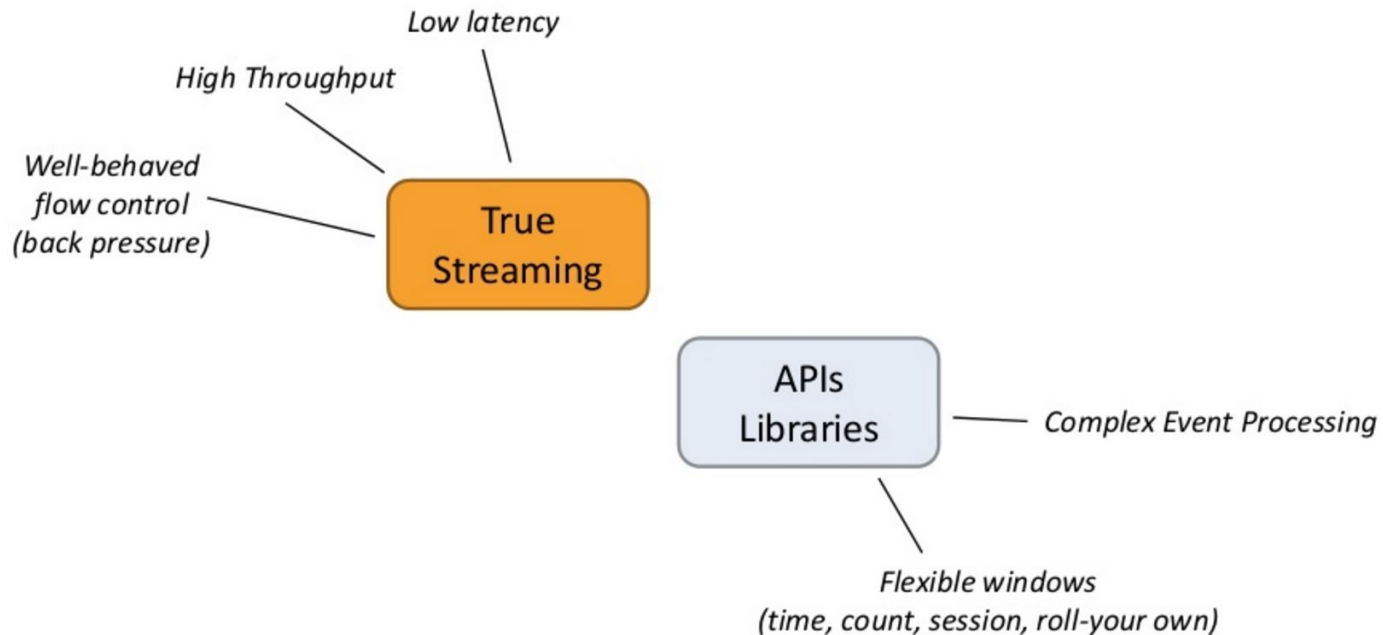
   @tzulitai

Flink.tw

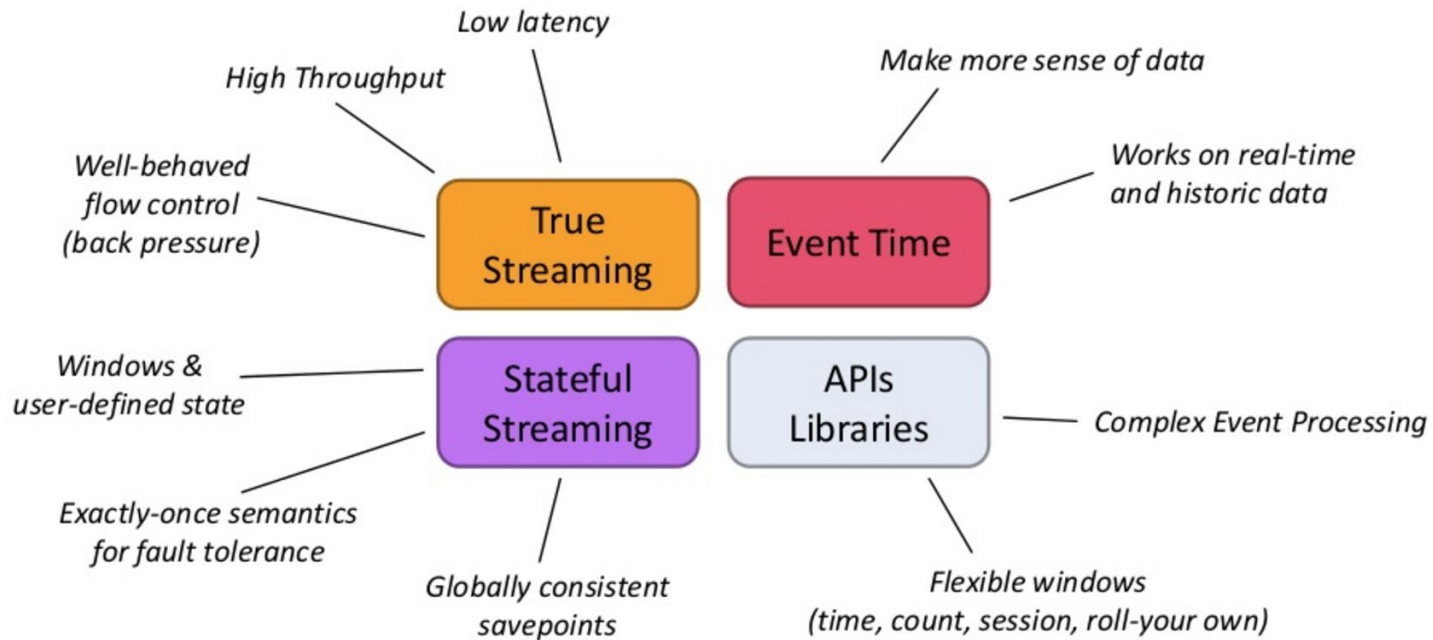
Apache Flink Taiwan User Group

Sept 2016 @ HadoopCon

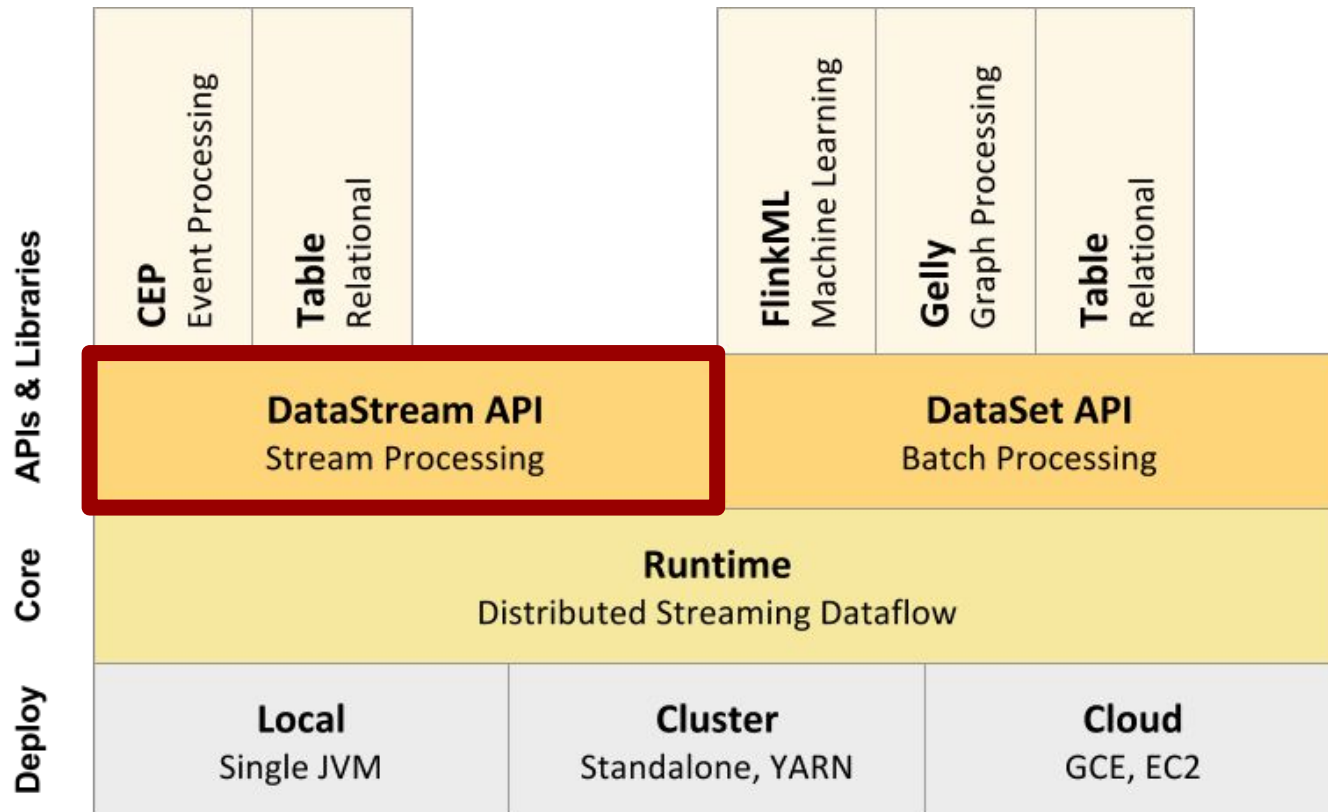
00 This session will be about ...



00 This session will be about ...



00 This session will be about ...



00 This session will be about ...

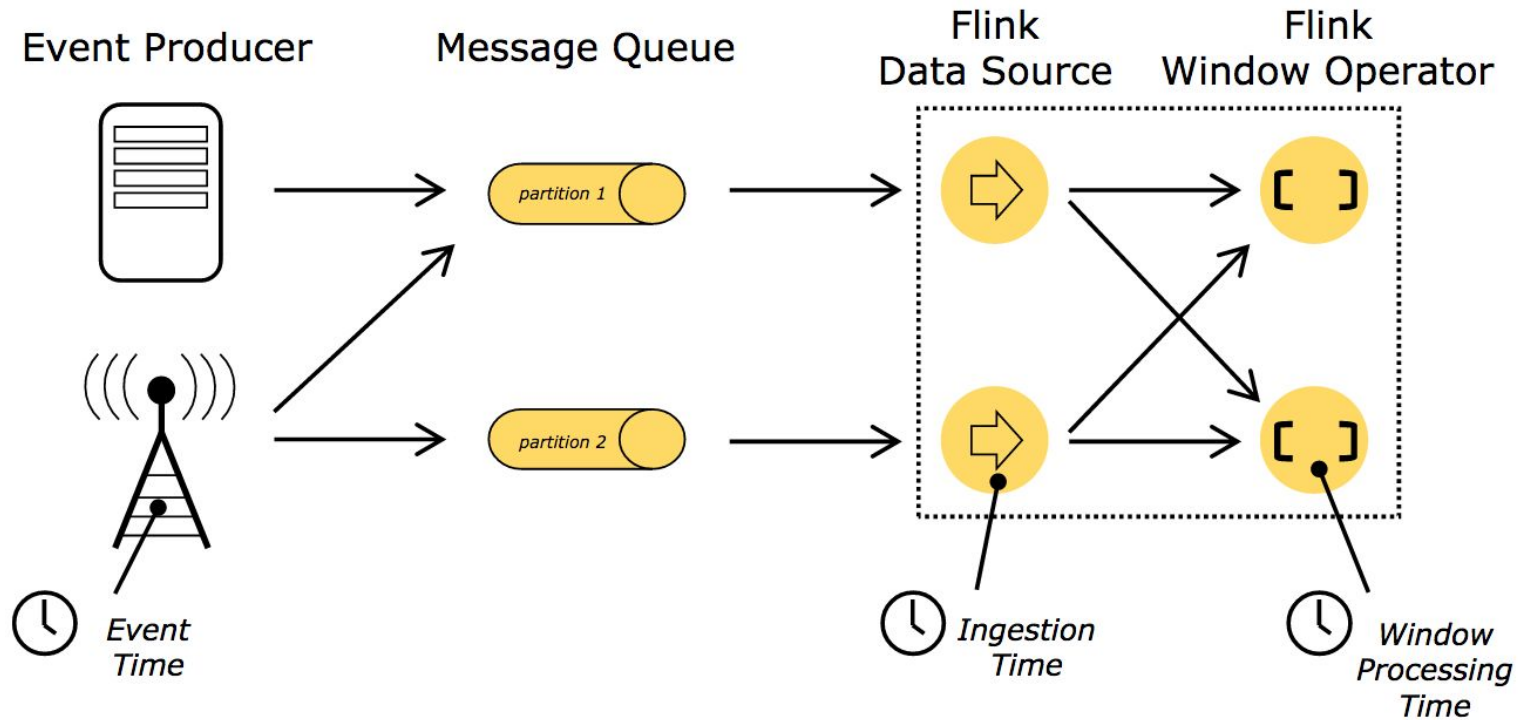
- Flink's notion of time in streaming jobs
- How *Watermarks* support Event-Time Processing
- Flink's fault-tolerant, exactly-once streaming semantics
- Flink's distributed snapshot checkpointing
- Out-of-core streaming state backends

Flink's Notions of Time

01 Different Kinds of “Time”

- ***Processing Time:***
 - The timestamp at which a system processes an event
 - “Wall Time”
- ***Ingestion Time:***
 - The timestamp at which a system receives an event
 - “Wall Time”
- ***Event Time:***
 - The timestamp at which an event is generated

01 Different Kinds of “Time”



02 Why Wall Time is Incorrect

- Think *Twitter hash-tag count every 5 minutes*
 - We would want the result to reflect the number of Twitter tweets actually tweeted in a 5 minute window
 - **Not** the number of tweet events the stream processor receives within 5 minutes

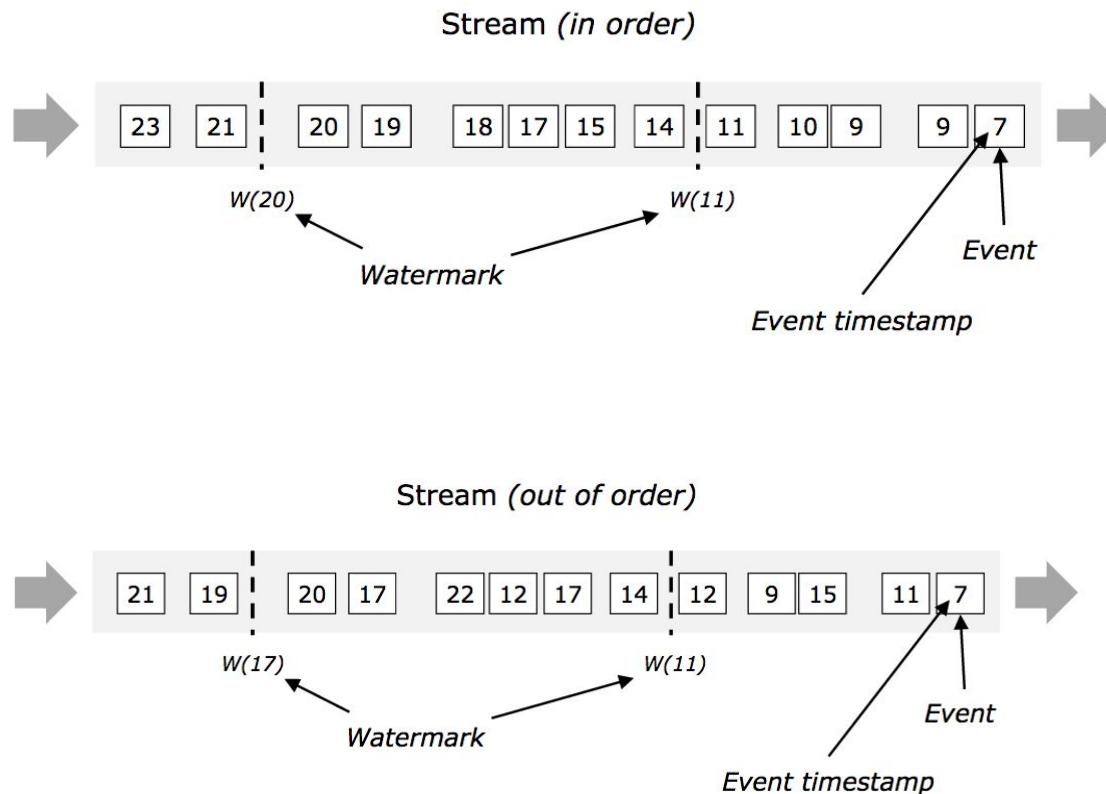
02 Why Wall Time is Incorrect

- Think *replaying a Kafka topic on a windowed streaming application ...*
 - If you're replaying a queue, windows are definitely wrong if using a wall clock

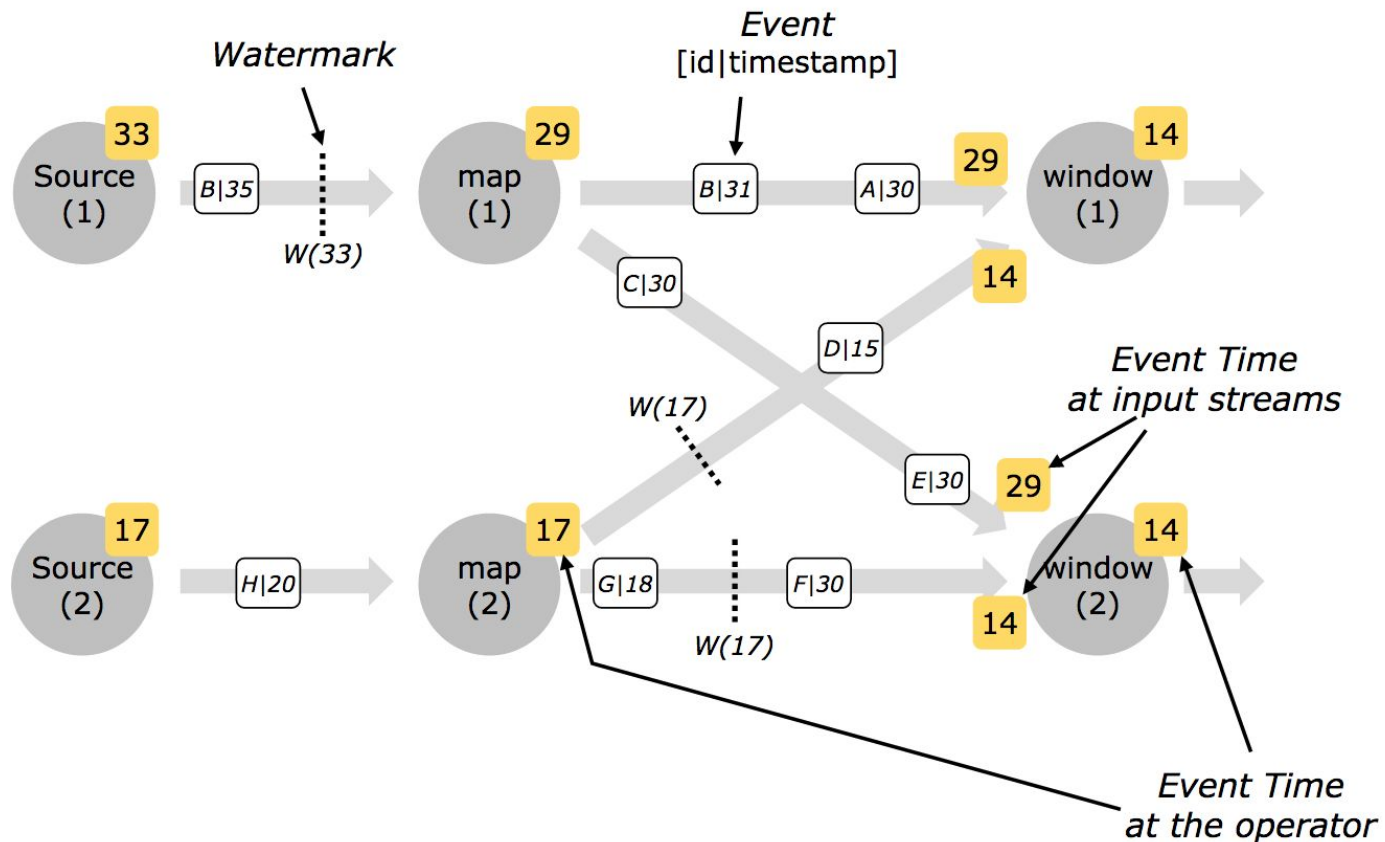
03 Watermarks & Event-Time

- **Watermarks** is a way to let Flink monitor the progress of event time
- Essentially a record that flows within the data stream
- Watermarks carry a timestamp t . When a task receives a t watermark, it knows that there will be no more events with timestamp $t' < t$

04 Watermarks & Event-Time



05 Watermarks in Parallel Streams



06 Event-Time Processing API

```
val env = StreamExecutionEnvironment.getExecutionEnvironment
env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime) ← Tell Flink to use “Event Time”
```

```
val env = StreamExecutionEnvironment.getExecutionEnvironment
env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime)

val stream: DataStream[MyEvent] = env.readFile(
    myFormat, myFilePath, FileProcessingMode.PROCESS_CONTINUOUSLY, 100,
    FilePATHFilter.createDefaultFilter());

val withTimestampsAndWatermarks: DataStream[MyEvent] = stream
    .filter( _.severity == WARNING )
    .assignTimestampsAndWatermarks(new MyTimestampsAndWatermarks())

withTimestampsAndWatermarks
    .keyBy( _.getGroup )
    .timeWindow(Time.seconds(10))
    .reduce( (a, b) => a.add(b) )
    .addSink(...)
```

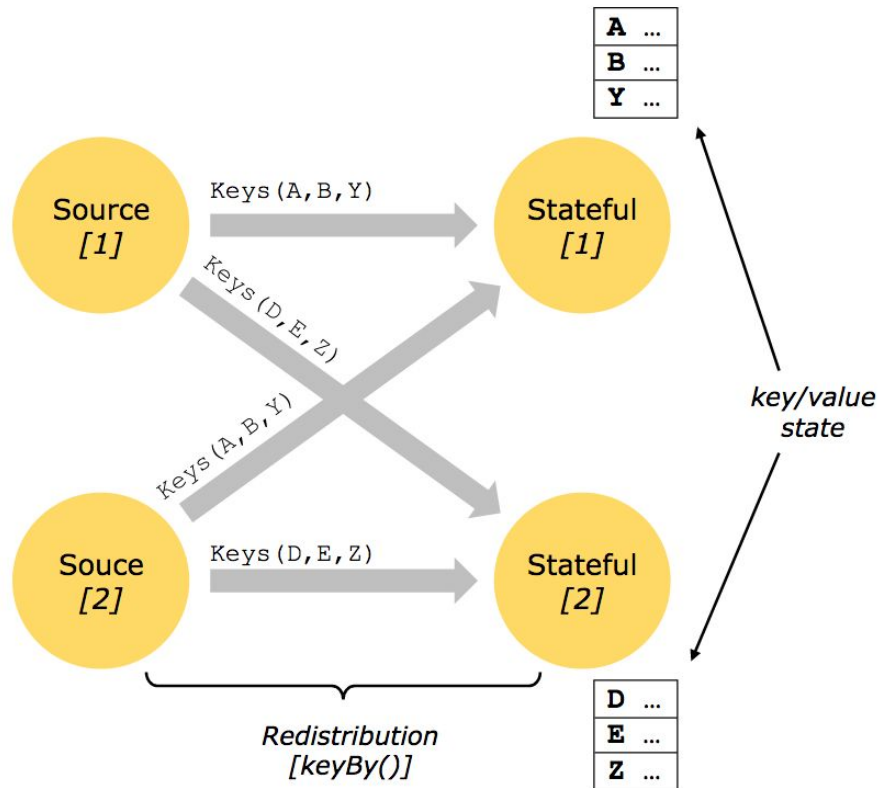
Assign event timestamps and watermarks

Exactly-Once Streaming Fault-Tolerance

07 Stateful Streaming

- **Any** non-trivial streaming application is stateful
- To draw insights from a stream you usually need to look beyond a single record
- Any kind of aggregation is stateful (ex. windows)


08 What “state” looks like in Flink



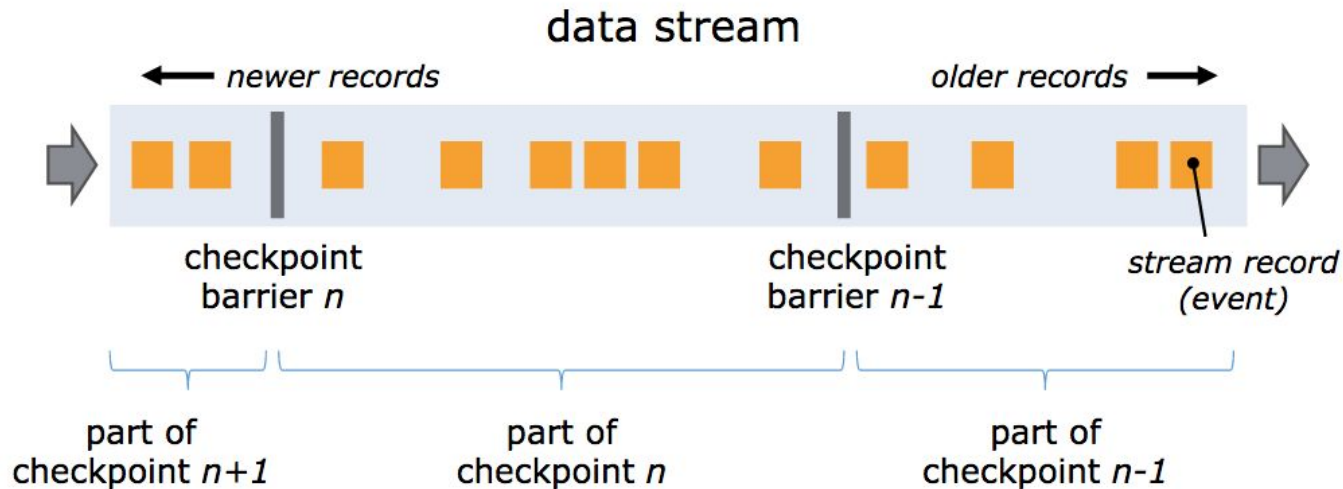
- Any Flink task can be stateful
- State is partitioned with the streams that are read by stateful tasks

09 Distributed Snapshots

- On each checkpoint trigger, task managers tell all stateful tasks that they manage to snapshot their own state
- When complete, send checkpoint acknowledgement to JobManager
- *Chandy Lamport Distributed Snapshot Algorithm*

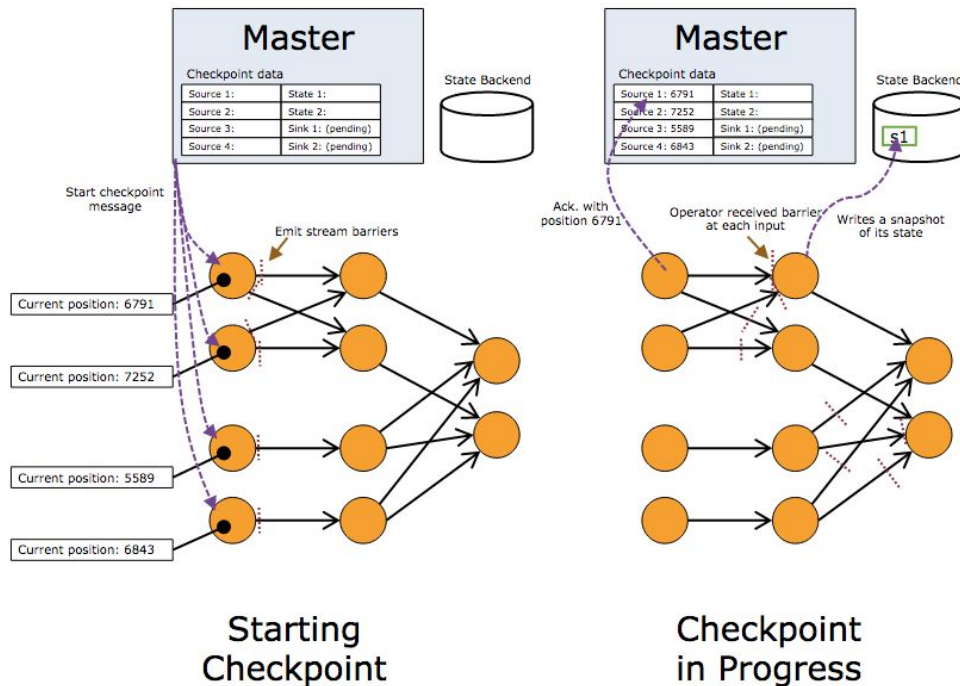


09 Distributed Snapshots



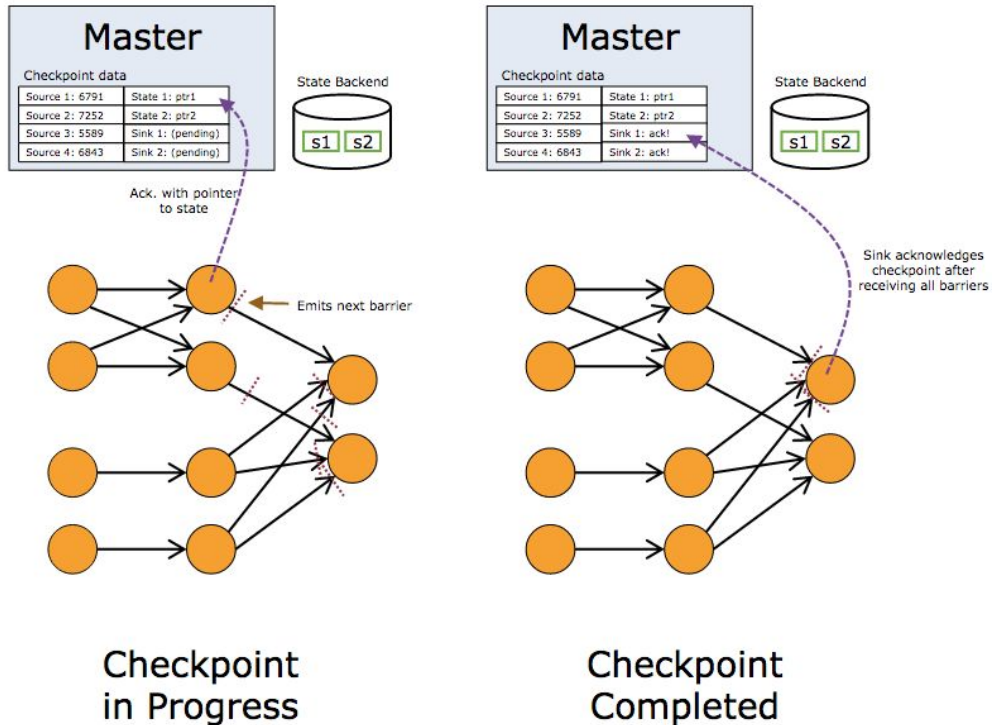
- On a checkpoint trigger by the JobManager, a ***checkpoint barrier*** is injected into the stream

10 Distributed Snapshots



- When a task receives a checkpoint barrier, its state is checkpointed to a state backend
- A pointer value to the stored state is stored in the distributed snapshot

11 Distributed Snapshots



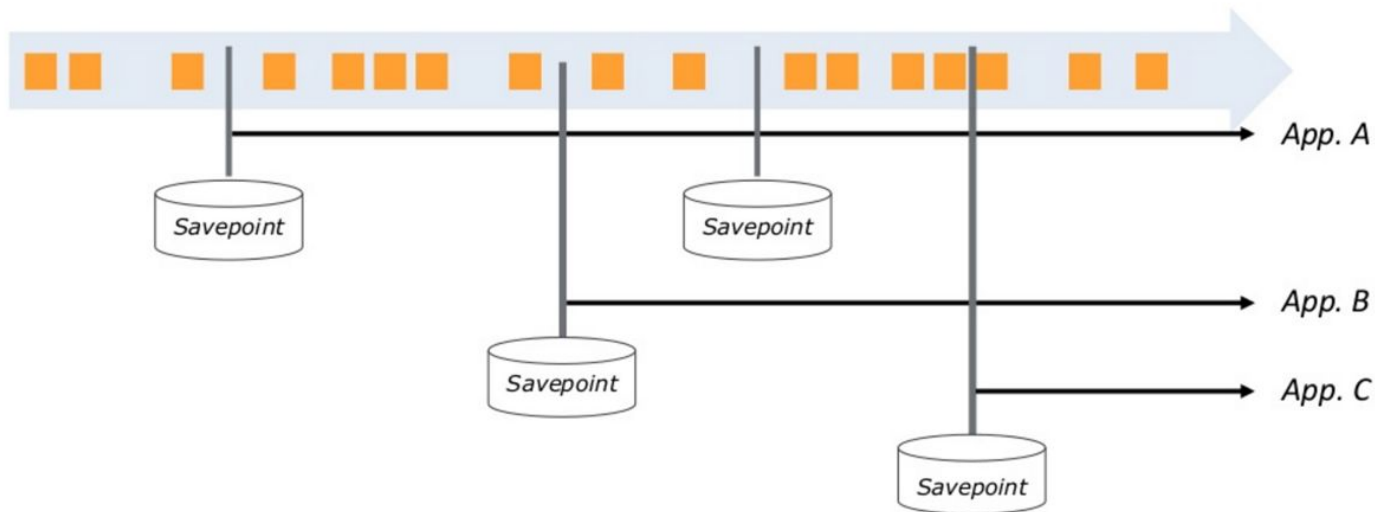
- After all stateful tasks acknowledges, the distributed snapshot is completed
- Only fully completed snapshots are used for restore on failure

12 Checkpointing API

```
val env = StreamExecutionEnvironment.getExecutionEnvironment()  
  
env.enableCheckpointing(100) // trigger checkpoint every 100ms  
env.setStateBackend(new RocksDBStateBackend(...))
```


13 Flink Streaming Savepoints

- Basically, a checkpointed that is persisted in the state backend
- Allows for stream progress “versioning”



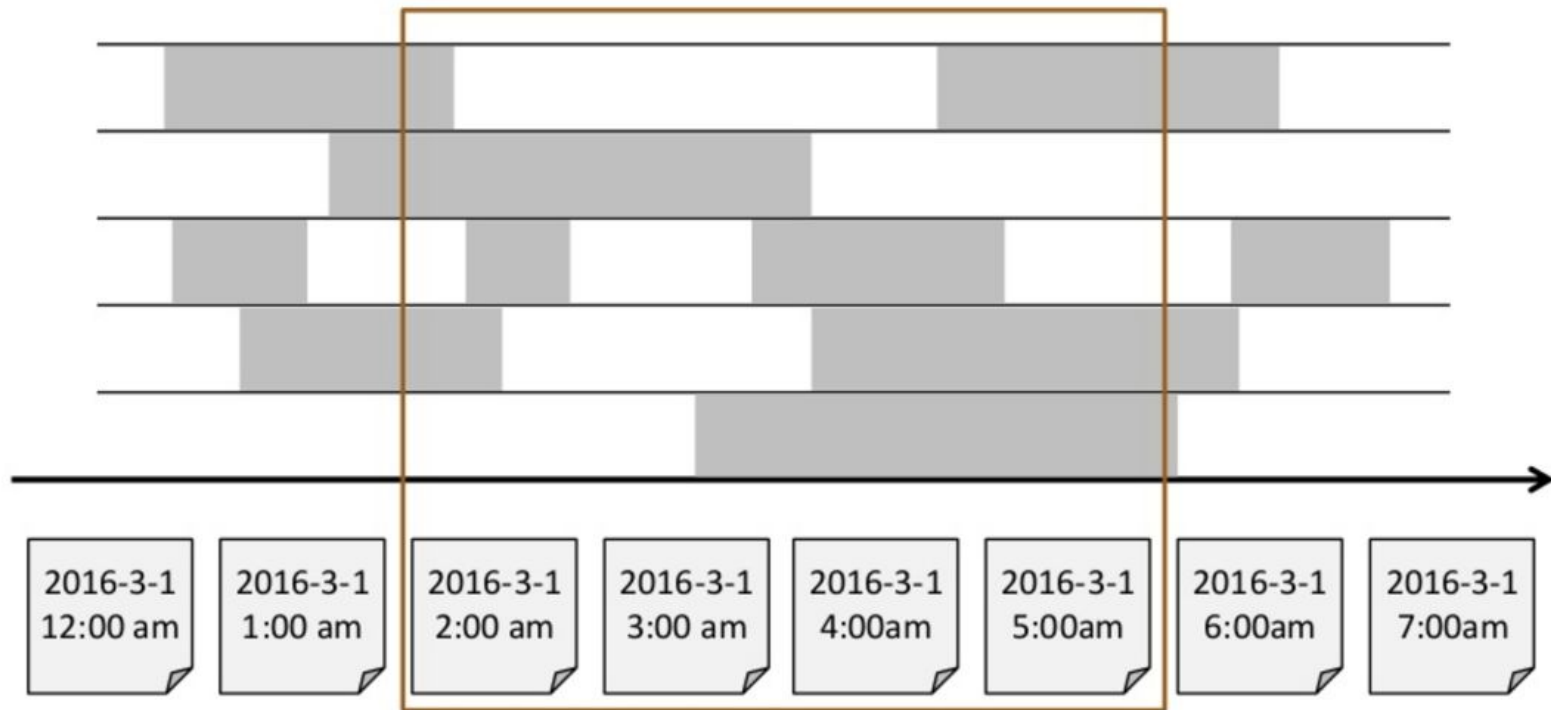
14 Power of Savepoints

Sessions over time




- No stateless point in time

14 Power of Savepoints




- Reprocessing as batch

14 Power of Savepoints




- Reprocessing as batch (*corrupt state*)

14 Power of Savepoints



- Reprocessing as streaming, starting from savepoint

15 Power of Savepoints



- Reprocessing as streaming, starting from savepoint