



Apache Flink™ Training

System Overview

Tzu-Li (Gordon) Tai

tzulitai@apache.org

   @tzulitai

Flink.tw

Apache Flink Taiwan User Group

Sept 2016 @ HadoopCon

00 Who am I?

- 戴資力 (Gordon)
- Apache Flink Committer
- Co-organizer of Apache Flink Taiwan User Group
- Software Engineer @ VMFive
- Java, Scala
- Enjoy developing distributed computing systems

00 This session will be about ...

- An understanding of Apache Flink
- Flink's streaming-first philosophy
- What exactly is a “streaming dataflow engine”?



Apache Flink

an open-source platform for distributed stream and batch data processing

- Apache Top-Level Project since Jan. 2015
- **Streaming Dataflow Engine** at its core
 - Low latency
 - High Throughput
 - Stateful
 - Distributed



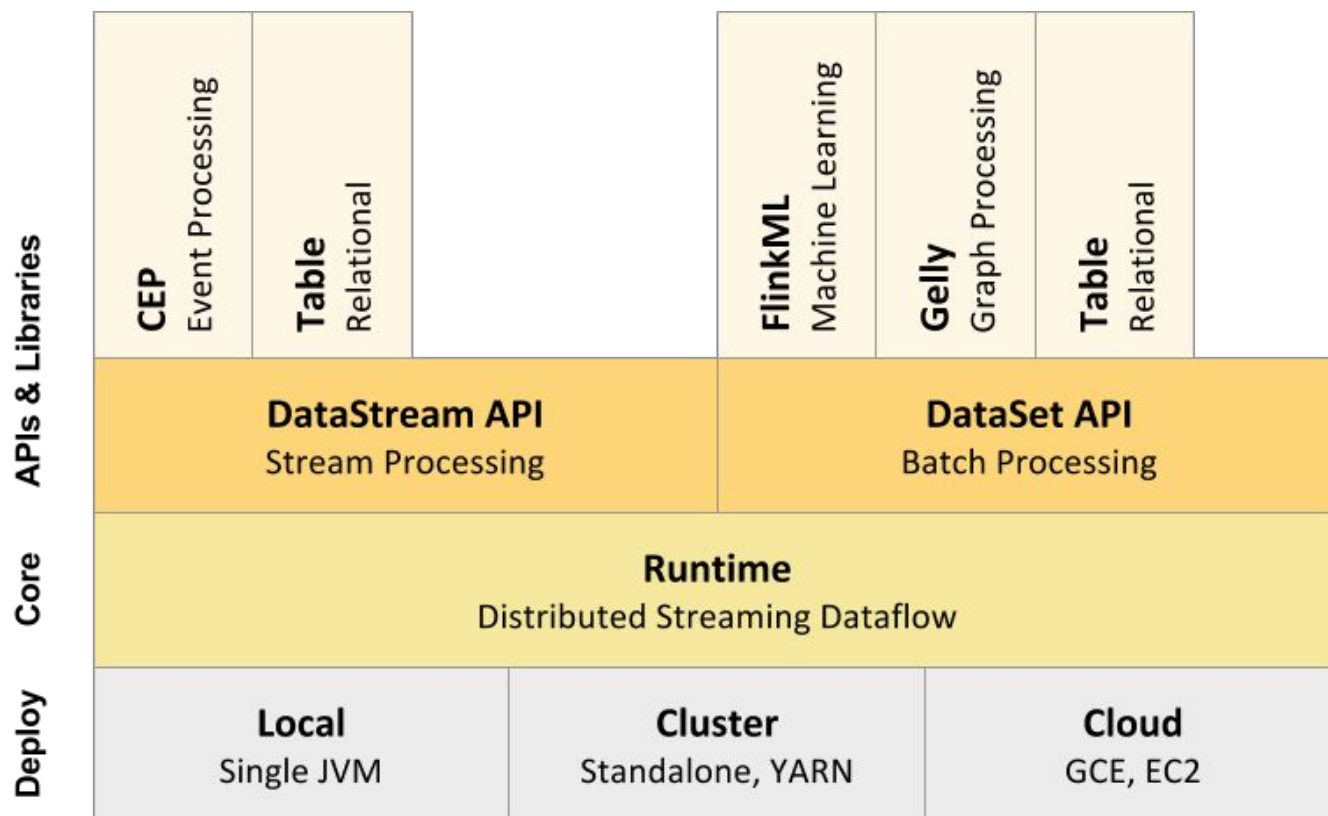


Apache Flink

an open-source platform for distributed stream and batch data processing

- ~100 active contributors for Flink 1.1.x release
- Used in **production** by:
 - **Alibaba** - realtime search ranking optimization
 - **Uber** - ride request fulfillment marketplace
 - **Netflix** - Stream Processing as a Service (SPaaS)
 - **Kings Gaming** - realtime data science dashboard
 - ...

01 Flink Components Stack



02 Scala Collection-like API

```
case class Word (word: String, count: Int)
```

DataSet API

```
val lines: DataSet[String] = env.readTextFile(...)

lines.flatMap(_.split(" ")).map(word => Word(word,1))
  .groupBy("word").sum("count")
  .print()
```

DataStream API

```
val lines: DataStream[String] = env.addSource(new KafkaSource(...))

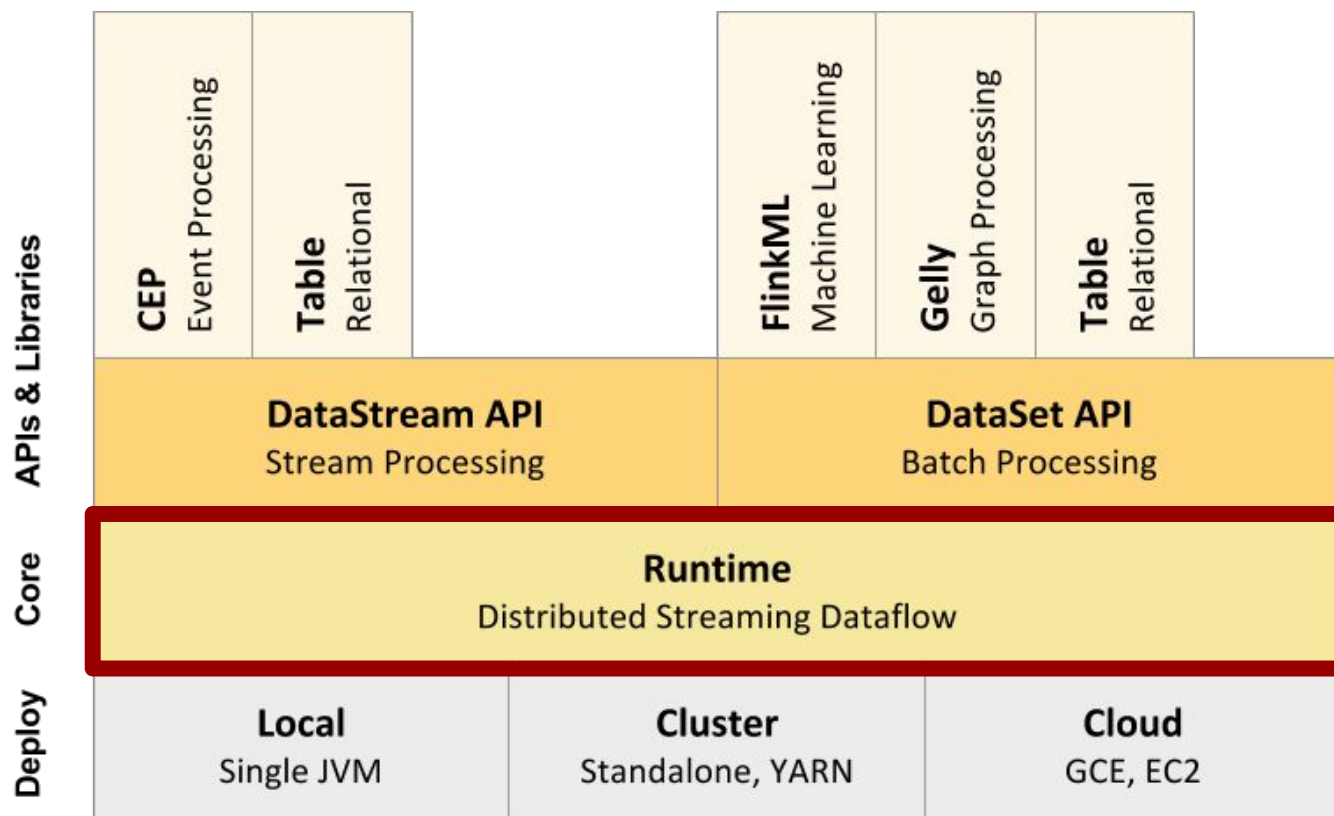
lines.flatMap(_.split(" ")).map(word => Word(word,1))
  .keyBy("word").timeWindow(Time.seconds(5)).sum("count")
  .print()
```

02 Scala Collection-like API

```
.filter(...).flatMap(...).map(...).groupBy(...).reduce(...)
```

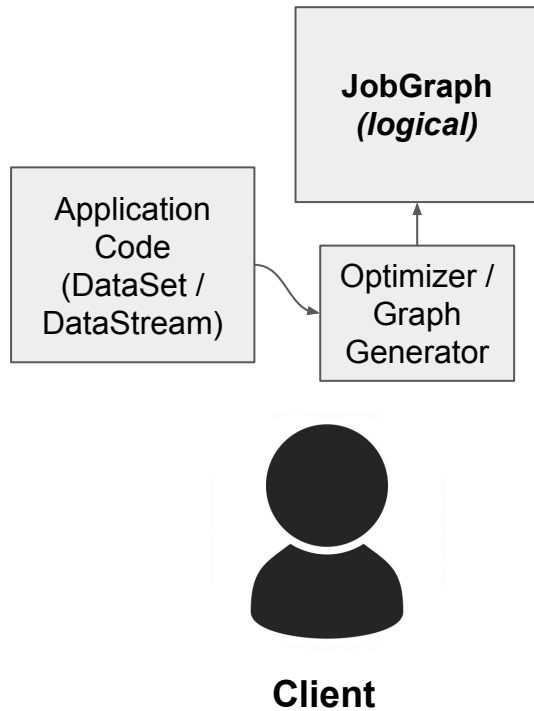
- Becoming the *de facto standard* for new generation API to express data pipelines
- Apache Spark, Apache Flink, Apache Beam ...

03 Focusing on the Engine ...



Lifetime of a Flink Job

04 Flink Job




04 Flink Job

Application code:

- Define sources
- Define transformations
- Define sinks

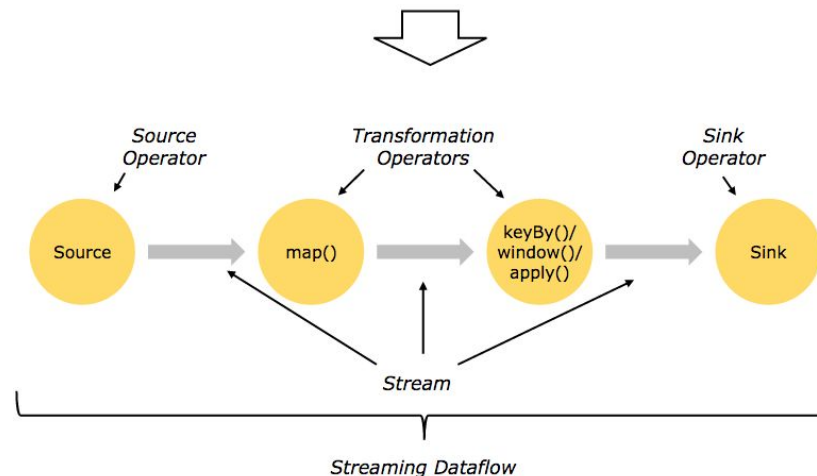
```
DataStream<String> lines = env.addSource(  
    new FlinkKafkaConsumer<> (...));  
DataStream<Event> events = lines.map((line) -> parse(line));  
DataStream<Statistics> stats = events  
    .keyBy("id")  
    .timeWindow(Time.seconds(10))  
    .apply(new MyWindowAggregationFunction());  
stats.addSink(new RollingSink(path));
```

Source
Transformation
Transformation
Sink

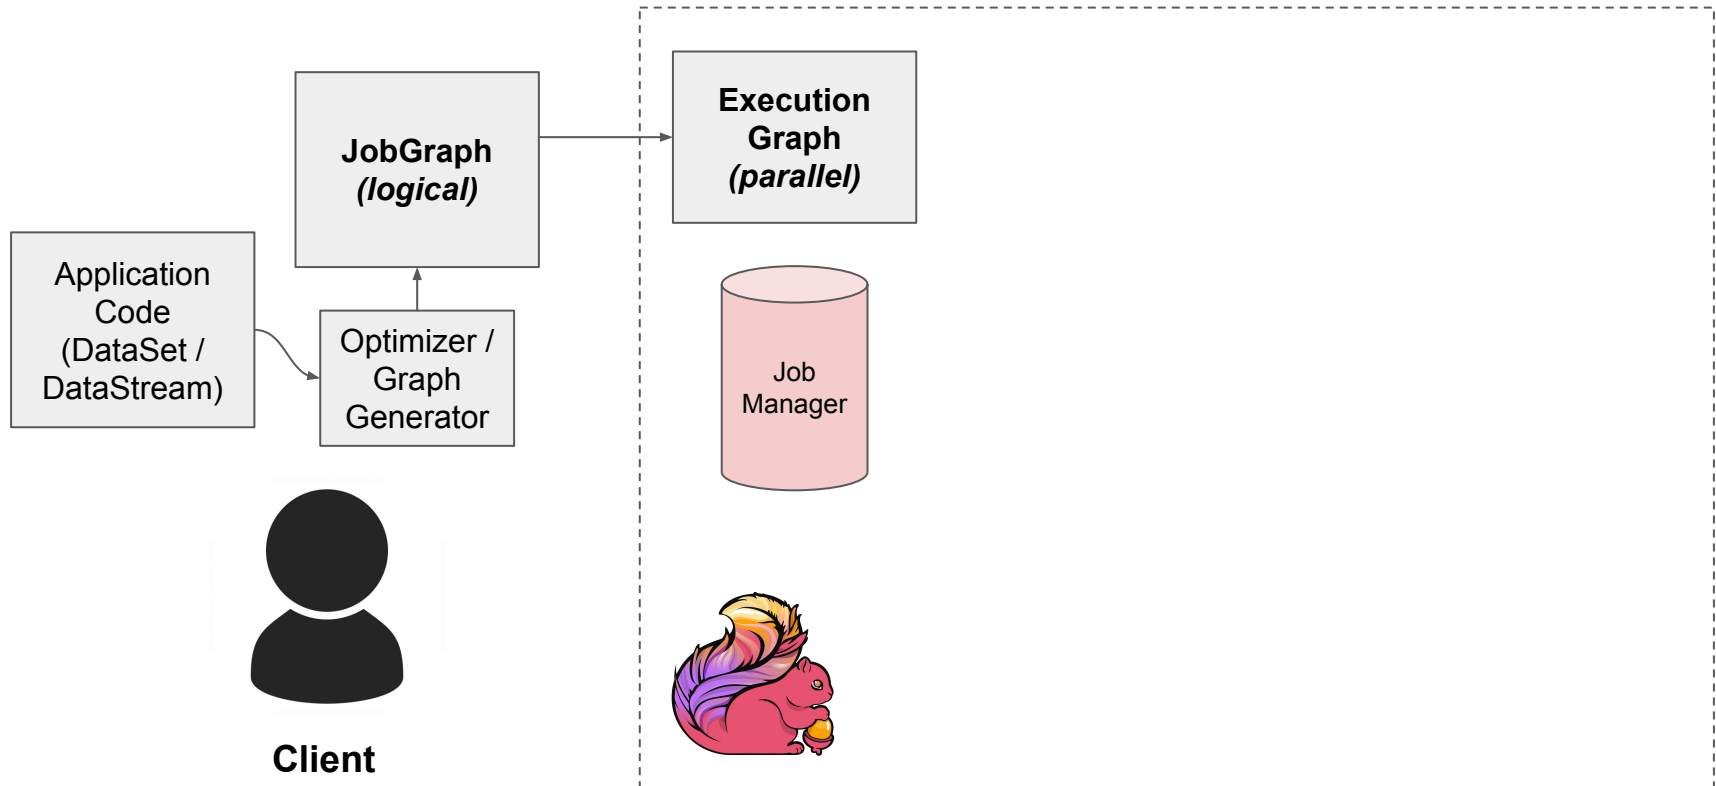


JobGraph

logical view of the
dataflow pipeline

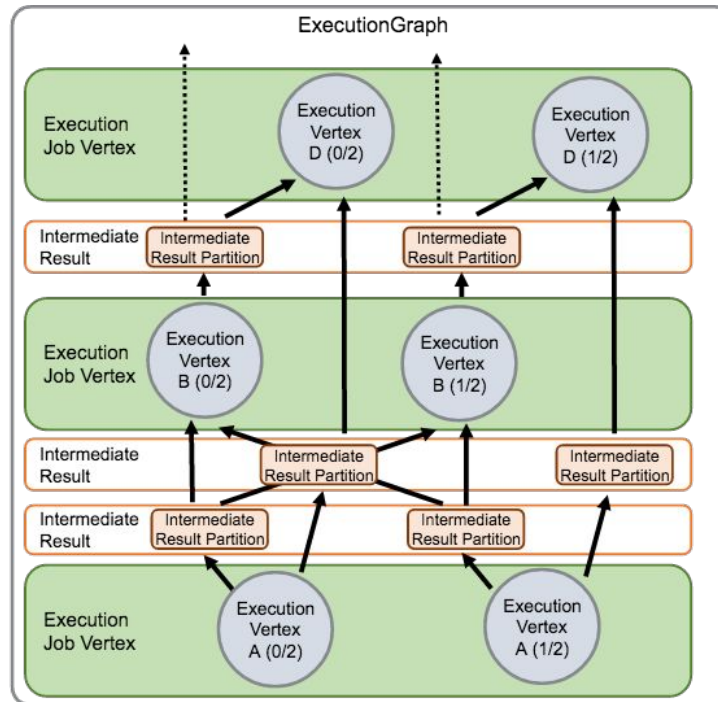
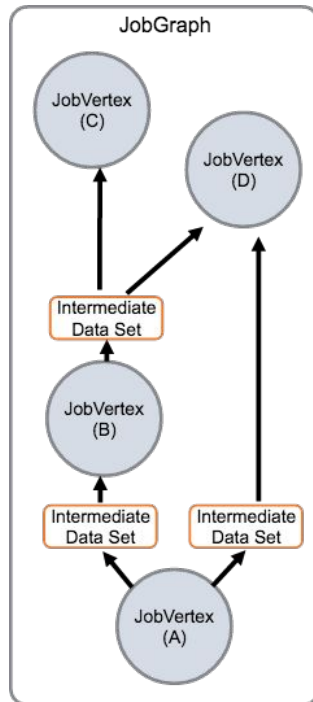


04 Flink Job



04 Flink Job


*logical
execution
plan*



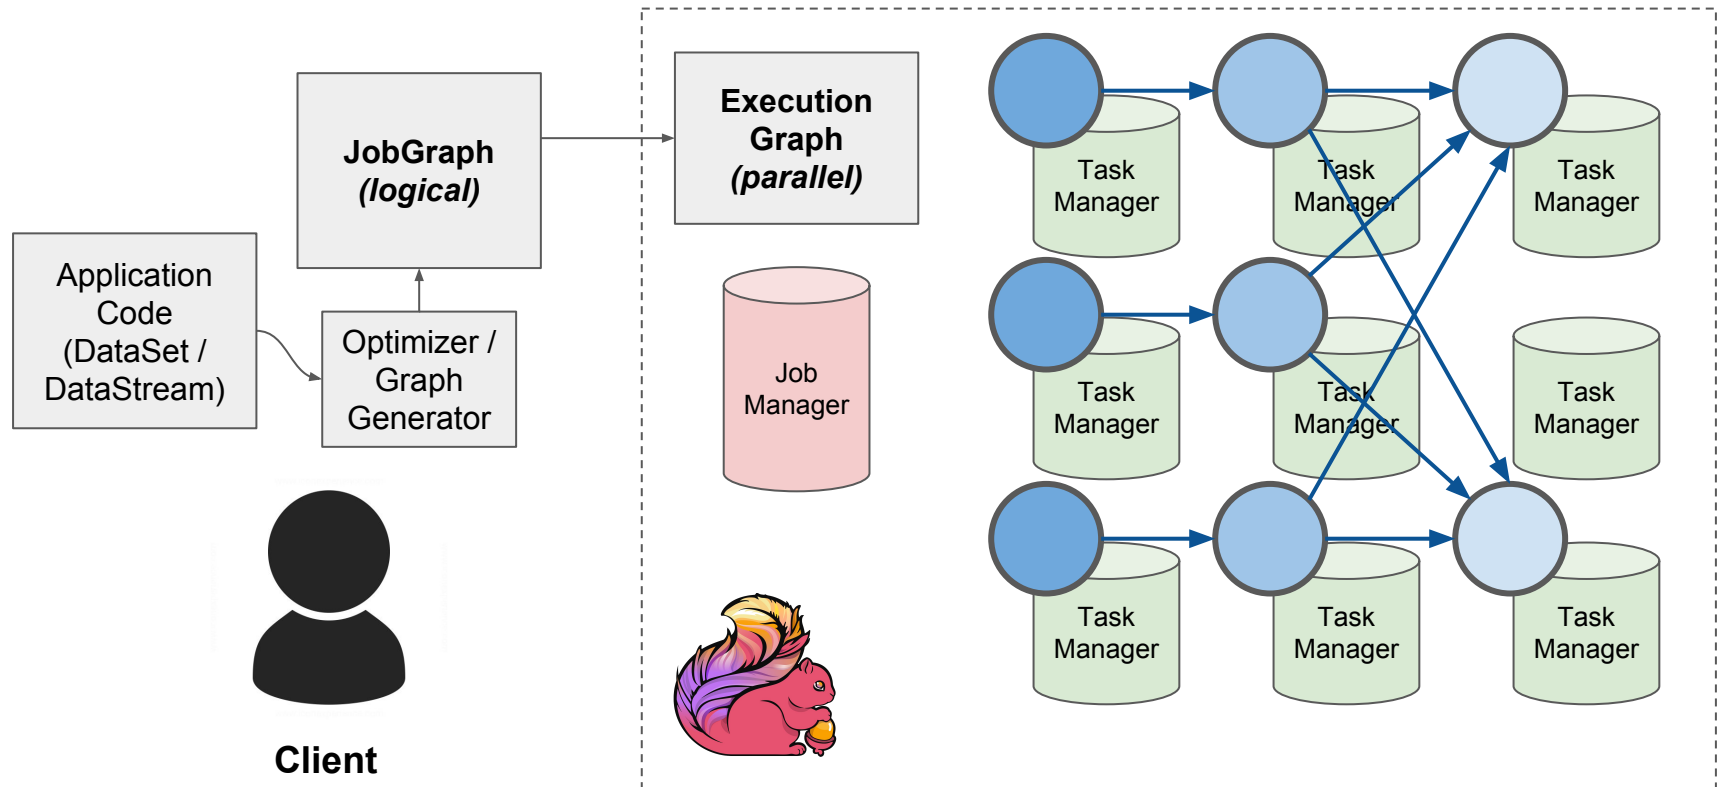
*physical
parallel
execution
plan*

→ **Intermediate Result Partitions** define a data shipping channel between parallel tasks (not materialized data in memory or disk)

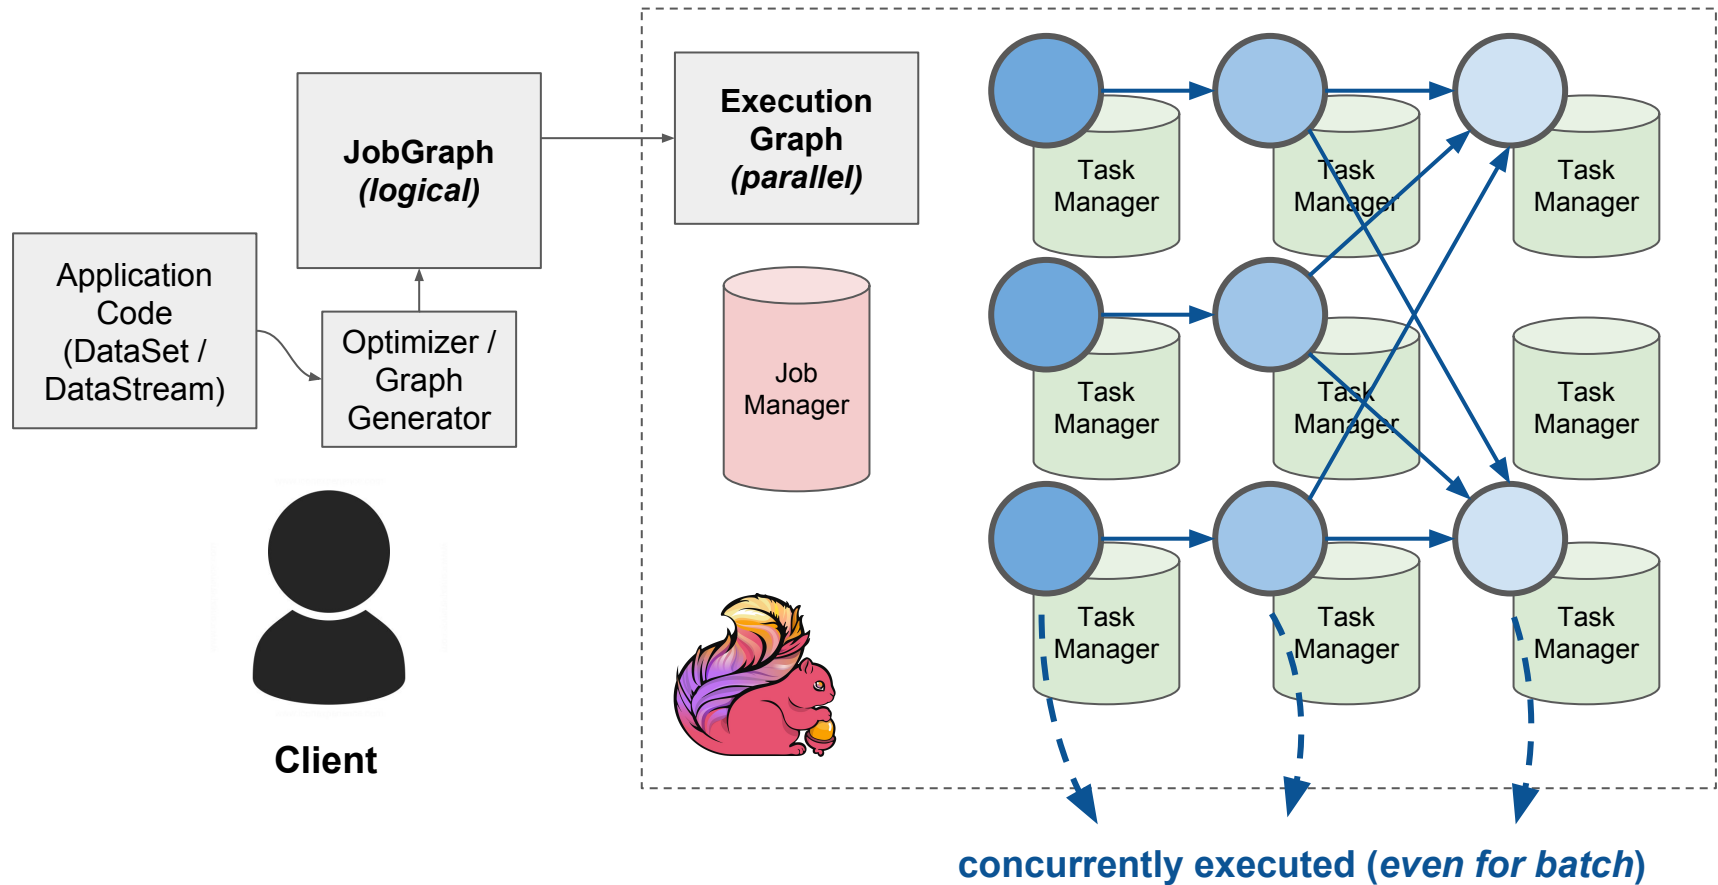
04 Flink Job



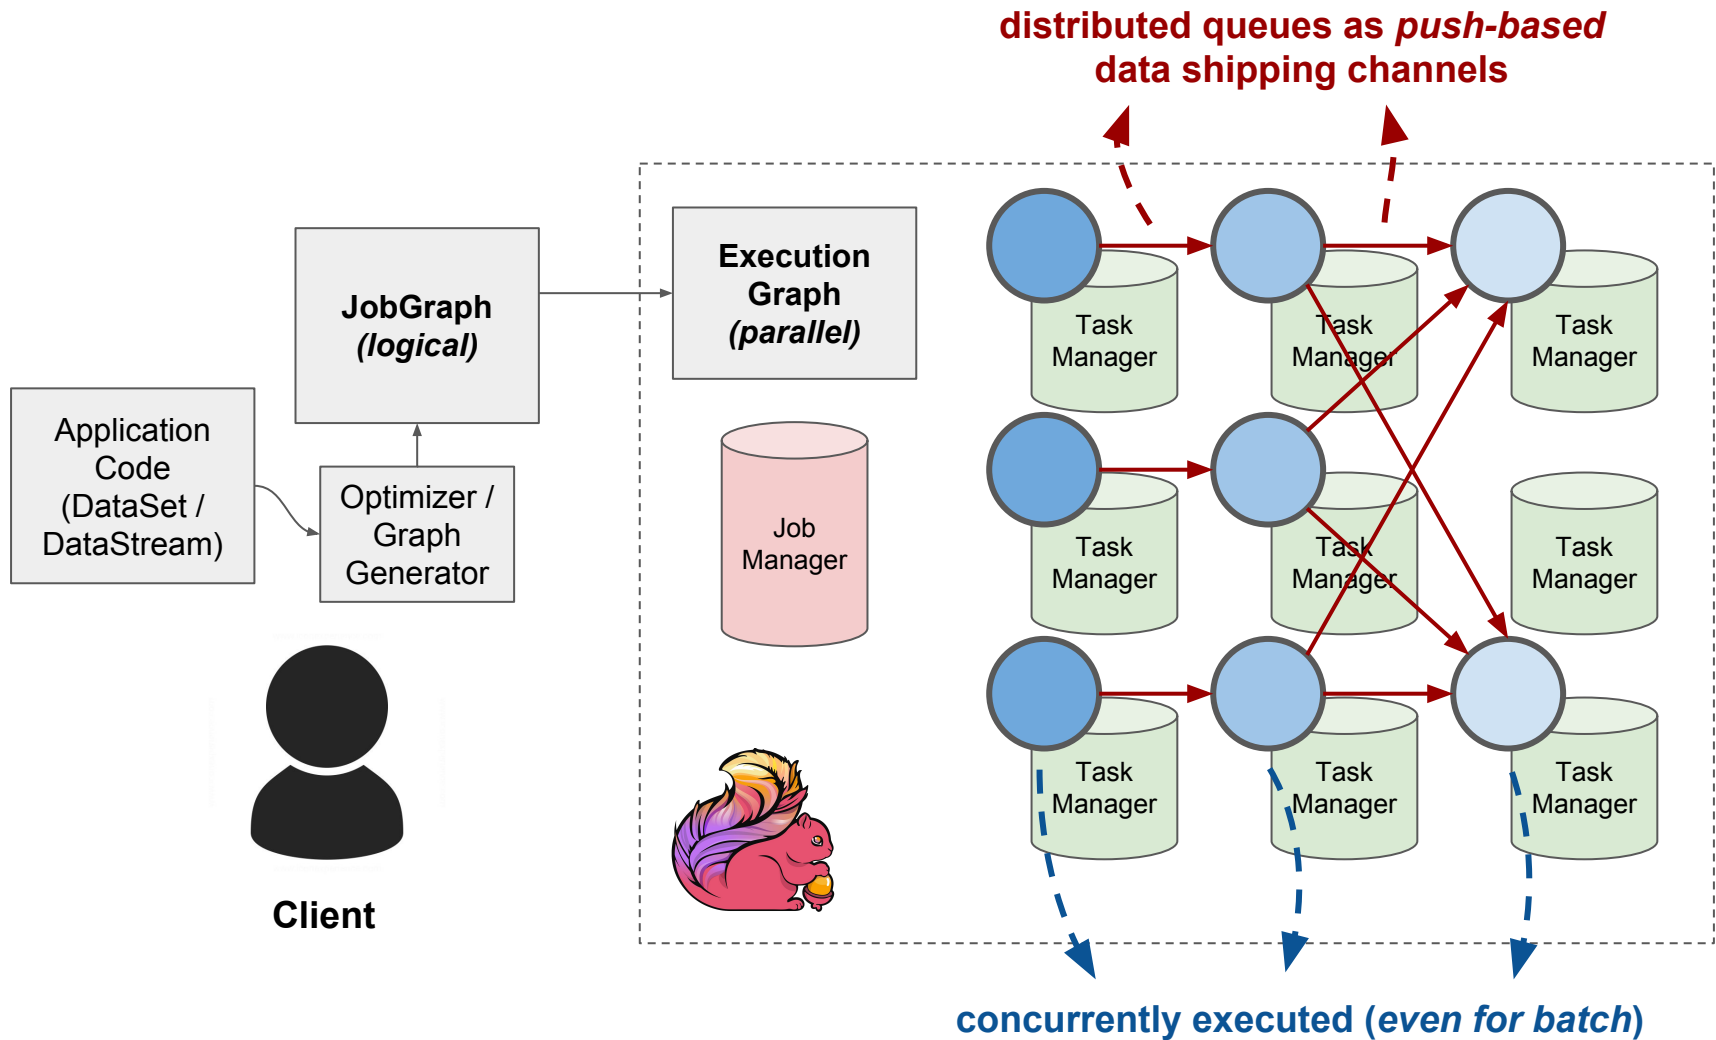
04 Flink Job



04 Flink Job




04 Flink Job



Advantages of a Streaming Dataflow Engine

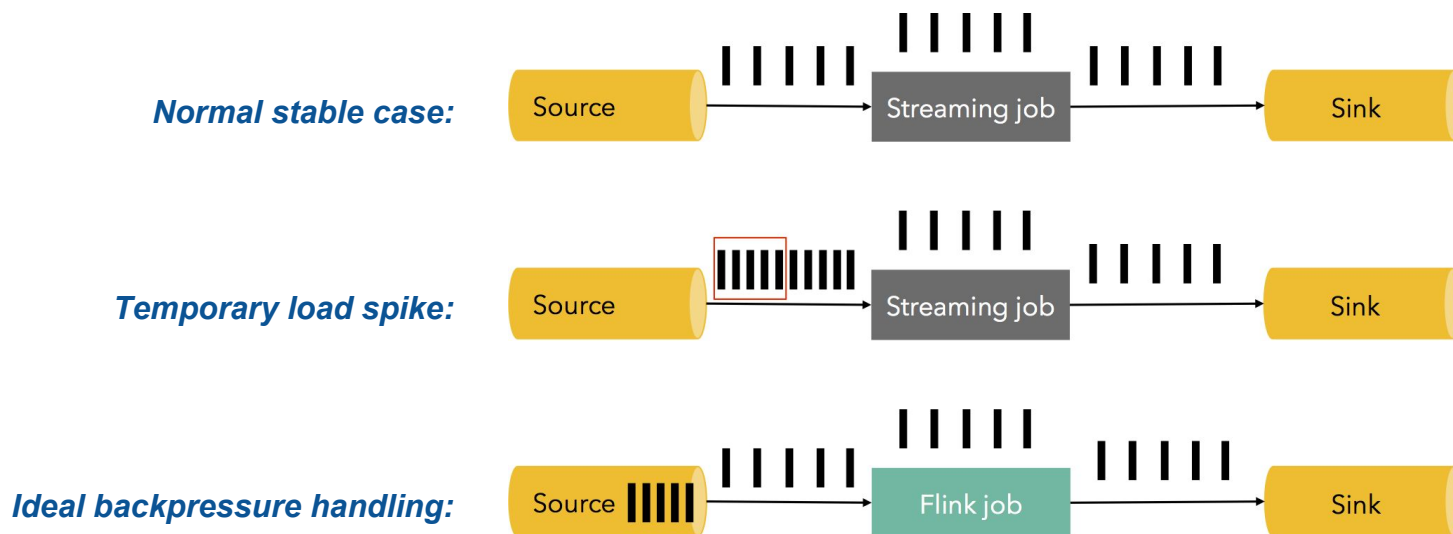
05 Streaming Dataflow Engine

- True one-at-a-time streaming
- Tasks are scheduled and executed concurrently
- Data transmission is *pushed* across *distributed queues*




06 Backpressure Handling

- Receiving data at a higher rate than a system can process during a temporary load spike
 - ex. GC pauses at processing tasks
 - ex. data source natural load spike




07 How Distributed Queues Work



1. Record "A" enters Task 1, and is processed
2. The record is serialized into an *output buffer* at Task 1
3. The buffer is shipped to Task 2's *input buffer*

Observation: Buffers need to be available throughout the process
(think *Java blocking queues used between threads*)


07 How Distributed Queues Work



1. Record "A" enters Task 1, and is processed *Taken from an output buffer pool*
2. The record is serialized into an **output buffer** at Task 1
3. The buffer is shipped to Task 2's **input buffer** *Taken from an input buffer pool*

Observation: Buffers need to be available throughout the process
(think Java blocking queues used between threads)

07 How Distributed Queues Work




- Size of buffer pools defines the amount of buffering Flink can do in presence of different sender / receiver speeds
- **More memory** simply means Flink can **simply buffer away** transient, short-living **backpressure**
- **Lesser memory** simply means more **immediate backpressure response**

08 Flexible Windows


- Due to one-at-a-time processing, Flink has very powerful built-in windowing (certainly among the best in the current streaming framework solutions)
 - Time-driven: *Tumbling window, Sliding window*
 - Data-driven: *Count window, Session window*

09 Time Windows

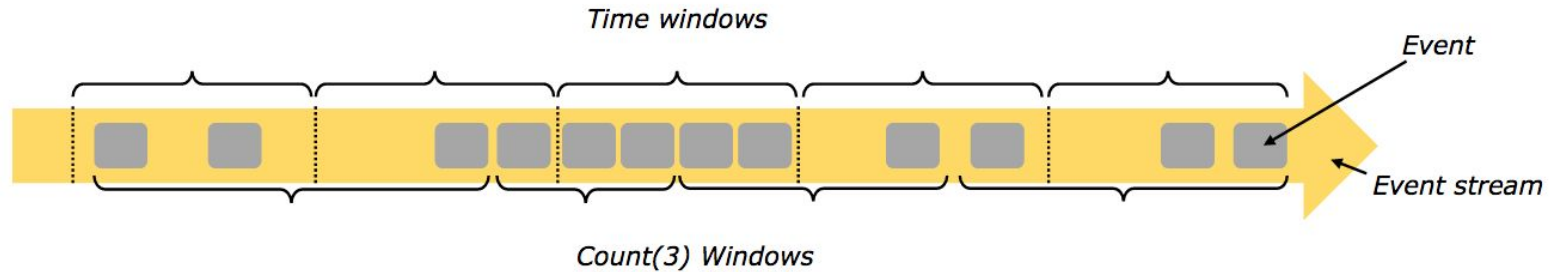
Tumbling Time Window




Sliding Time Window



10 Count-Triggered Windows



11 Session Windows



Prepare for Hands-On!

12 Clone & Import Project

→ <https://github.com/flink-taiwan/hadoopcon2016-training>